

Final Report

Smart Systems Lab: 360-degree camera

de Guzman, Carlos
Ng, Alexander

Früge, Alan
Park, Daniel

Kirby, William

I. Abstract

With the advancement of computer vision, the applications of cameras have dramatically increased. In this project, the goal is to create a 360° camera using an Nvidia Jetson Nano stitching several camera frames. To achieve this goal, a custom case and stitching software were created for this project. The algorithm for the stitching software will be improved iteratively and is the greatest challenge in creating a reliable camera system. Also included is a basic and general-purpose GUI that communicates with the Jetson from a host computer.

This report begins with how the general-purpose framework of this project could be reused for a wide variety of applications. Next, it will provide the technical background regarding computer vision as well as our test results with different camera and stitcher combinations. These informed the decision to create a custom stitching algorithm. The implementation of the custom stitcher is documented in the design section. This section will also include the communication protocols that interface between each of the subsystems and how the overall system should interact from both a user and developer perspective. Included is a timeline of the completed features. Based on the limitations of the completed features, more specific recommendations for the use cases are discussed in the impact section. Lastly, a brief discussion of future modifications and applications will be discussed.

II. Introduction

This project aims to create an inexpensive and reliable 360° camera for both consumers and businesses. As the completed product will be able to capture an entire scene around it, it has a wide variety of applications. Examples include panoramic video conferencing, virtual experience capturing for VR, etc.

While there are currently existing 360° cameras available to consumers, it is important to note that the software and design artifacts are made open source on Github. This will allow an end user to build upon the general-purpose architecture to fulfill the requirements of a specific application. Also, documented there is the specific hardware that was used. Because the software was built on the UNIX platform, an end user could also choose to change the hardware to better fit a use case. An example of possible changes could include using a Raspberry Pi over the NVIDIA Jetson Nano, different cameras, and redesigning a case to include more cameras.

An example application proposed in the SmartSystems Lab was the use of computer vision to improve businesses' efficiency with tasks like inventory and security statistics. In this case, the user could experiment with different camera parameters including the number of connected cameras or the stitcher being used. Another possibility is the use of an embedded system with higher computational power that could be used to run a lightweight neural network for item detection. Regardless of the specific changes, the goal of the project is to have a well-documented and general-purpose platform that is easy to build upon for different applications.

III. Background

A. Camera Calibration

As computer vision is the field within artificial intelligence that studies how computers interpret and understand the visual world, this paper will begin with cameras. Camera calibration is a foundational concept that determines how a camera represents an external scene as an image based on the extrinsic and intrinsic parameters of the camera [14]. Extrinsic parameters correspond to rotation and translation vectors from 3D points in the real world to the camera coordinate system. Conversely, intrinsic parameters convert the internal coordinate representation to the image plane. These include the focal length and the optical center of the camera which are unique to a camera based on its physical geometry, sensor, etc.

To compute these matrices, several images of a well-defined pattern like a chessboard with known real-world lengths should be taken. From these, OpenCV can interpolate the values of these matrices and save them. With these values known, it is possible to correct distortions which are important for stereo applications such as this project. In this project, it is also explored how intrinsic matrices are used for cylindrical warping explained below.

$$\text{camera matrix} = \begin{bmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

Figure 0: General Camera matrix (focal length and optical center) [14]

B. Homography

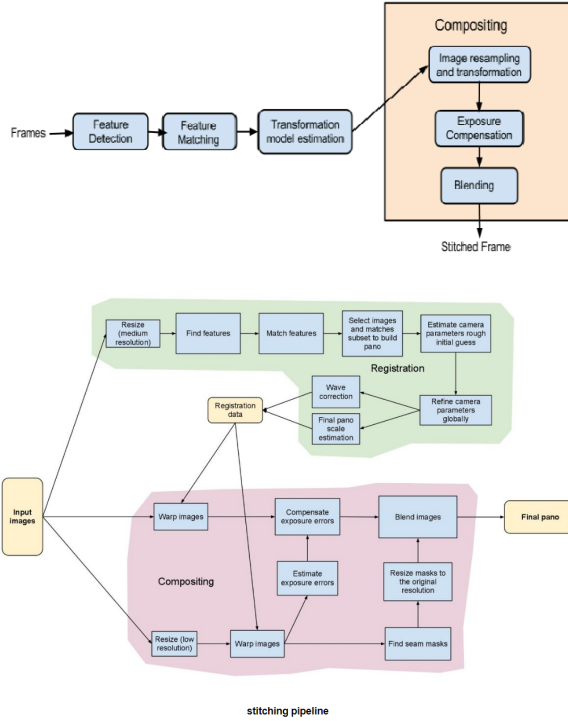


Figure 1: Image Stitching Pipelines

In the generalized image stitching pipeline, shown above, the stages can be grouped into transformation estimation (first three) and compositing (final three). One method of the former is using homography matrices that relate [15]. It is composed of a set of rotations, translations, and plane normal matrices. By extracting and matching features between 2 images, it is possible to estimate this homography matrix. Using a homography matrix, it is possible to warp the perspective of one image and overlay it over the other. This is shown in the 2 figures below.




Figure 2: Homography with matched features [15]




Figure 3: Planar projection/warping

C. Compositing and Cylindrical Warping

Compositing applies the estimated transforms on the original images to create the final stitched image. The first step involves warping the images. Then exposure compensation and blending can be applied to improve the quality of the stitch. These two will not be discussed as complex techniques like exposure estimation and finding the best seam masks are beyond the scope of this project.

Cylindrical warping can be used for image compositing to project all images onto a common cylindrical plane. It requires the camera's intrinsic matrix and rotation matrices to be known. The image can then be “unrolled” to be equirectangular. While this may cause some distortion, it will allow the stitched image to span greater than 180 degrees. This is ideal for cases where there is one rotating camera or multiple cameras with a common center.




Figure 4: Cylindrical projection/warping

D. Testing




Figure 5: Runtimes of Image Stitching Stages [10]

Above is a chart from a study from Harvard University regarding Real-Time image stitching. It is important to notice that transformation estimation (keypoint/feature detection, description, matching) takes a significant majority of the execution time when compared with compositing (warping and stitching). This informed the decision to create a custom stitcher for this project. A testing methodology for the stitcher was created to ensure that optimizations improved the performance over the baseline (OpenCV) stitcher.




Figure 6: Stitch test results with OpenCV Stitcher




Figure 7: Stitch test results with Custom Stitcher

The key results from testing were as follows. As desired, the custom stitcher had a nearly 20x faster execution time when compared with the OpenCV stitcher. It also had an effective 100% stitch rate compared with only 17% for OpenCV. It is also noteworthy that the execution time of the custom stitcher did not have outliers like those of its counterpart. This contributed to a more consistent experience.

IV. Timeline

Pre-alpha Build (October 2nd)

- 3D printed prototype case created
 - Holds 7 cameras at a fixed angle
- Concurrently read from 4 cameras
 - Video artifacts and occasional crashing
- Utilized OpenCV stitcher to successively stitch 2 frames into video stream
 - Resized captured frames in effort to reduce stitch time
- Displayed performance metrics based on stitched frames
 - Low success rate
 - Low framerate (high stitch time)

Alpha Build (September 9th)

- Concurrently read from 6 cameras connected to the Jetson
 - Tested and found that multithreading did not solve issue
 - USB bandwidth issue resolved by limiting resolution and framerate at a driver level (instead of resizing)
 - USB power requirements of cameras fulfilled with externally powered HUB
- Alternate GUI program
 - Sends SSH commands from the host to Jetson
 - Utilizes socket to receive video data

Beta Build (September (28th)

- Custom Stitching Algorithm
 - Able to save homography matrix for sets of 3-camera inputs
 - Symlinked USB cameras according to position in case
 - ensures saved homography remains applicable between when it was computed and will be used
- Testing Program
 - Able to run for a given duration
 - Able to test the different camera and stitcher configurations
 - Creates detailed log files
 - Takes an image before running the test
 - Detailed testing results analyzer
 - Creates report with statistics and charts from logs and images
- Flask Server
 - Initializes Camera System
 - Recalibrates to environment
 - Routes created to send video through HTTP responses

Release Candidate Check-in (October 19th)

- Updated stitcher to stitch 2 sets of 3 images together
 - Noticeable seam between the 2 sets of images
- Updated GUI to be Web-based for improved expandability

- Updated Flask routes to properly serve the GUI

Release Candidate (October 28th)

- Improve the non-developer use case
 - Flask Runs on startup
 - Power Switch
- Cylindrical Warping
 - Manually tested different focal lengths effect on stitch quality
- Integrated *Panellum* Web Plugin video viewer (GUI)
 - Utilizes lightweight WebGL to dynamically change the perspective of a stitched image
 - Chosen because it is lightweight enough to run browser/client side without increasing the Jetson's load

Production Release (November 29th)

- Finalized 3D-printed shell accommodating all desired hardware including cameras, hub, LED indicators, and Jetson
- CUDA Acceleration
 - Compiled OpenCV library for CUDA
 - Verified install on the correct environment
- Improved camera calibration and matrix adjustments for cylindrical warping
- GUI improvements
 - Improved spacing
 - Updated buttons to fit with the Panellum framework
 - Added directions to improve user experience

V. Design

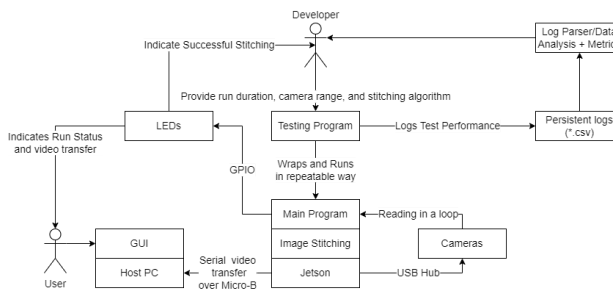


Figure 8: Behavior Diagram

This project was developed to run on the Nvidia Jetson Nano. This was chosen due to the CUDA acceleration which can help with many vectorized operations in computer vision and opens the possibility for expansions on this project like deployed image detection, etc. The Jetson connects to 6 cameras through an externally powered USB hub. These cameras are held in place with 60 degrees rotation along with the Jetson in a 3D printed case. The case was also designed to hold an LED indicator.

As discussed above, it was decided that a custom stitcher should be used. This works by taking images, estimating the transform, and then saving. It then uses this saved

transformation estimation to composited all subsequent image frames, so the output appears as a unified smooth video. Overall, this option was chosen as it was decided that with the computation constraints, fast execution and consistency were more important to the end user experience than perfect image quality.

The Jetson was configured to create a Flask web server on startup. This sets up the cameras and image stitcher and has routes to read stitched video. The Jetson can be connected to a host computer through serial Micro USB. This allows the host computer to communicate over IP. There is also an LED indicator for when the Flask server is running.

From the host computer, the user can open a web-based GUI. This program interacts with Jetson's flask web server through HTML requests. The interface includes buttons to refresh the video feed and recalibrate the cameras by estimating the transformation of images with the current environment (the cameras calibrate on start-up). Currently, it is being updated to use Pannellum which uses WebGL to allow the user to pan around the video. In addition, the user will have access to LEDs that indicate the stitching software's status.

On the developer side, the project includes a testing tool for different configurations in a repeatable way. This application takes a stitcher, execution time, and the number of cameras as parameters and runs the test while logging the results. From these results, a report can be generated. This includes graphs about stitch timing like figures 6 and 7 above, distribution of stitch errors along with a sample image to visually show the quality of the stitching.



Figure 9: Sample image with cylindrical warping (production release)

VI. Results

The camera system and all basic functions are complete with the implementation of a visually appealing user interface. In addition, the camera system case's redesign was complete and implemented to accommodate the Jetson and other hardware. Although the stitching algorithm has been iteratively improved and now contains at least one step from each of the stages in the image stitching pipeline, more options like those not included in the OpenCV library were not fully explored due to time constraints.

When choosing an output resolution, quality was sacrificed for performance. While the system is capable of outputting the stitched video at 480x640, it was decided that the time trade-off was too costly. Especially because of the time required for stitching and homography calculations.

In addition to the performance testing results above, end-to-end testing was performed with different stitching algorithms. Specifically, the impact of CUDA acceleration, camera calibration, and blending on framerate and stitch quality were analyzed. While it didn't have a significant performance impact, CUDA acceleration allowed for six camera streams to be stitched real time. Camera calibration did not have an effect on the stitch time or quality, so the cylindrical warp was manually tuned with different parameters to find what worked best with our specific camera setup. Lastly, blending was explored and found to have a noticeable impact on fps. This is why it was given as an option on the GUI to allow the user to choose whether to prioritize fps or quality.

This project was also made to be able to be used by non-developers. From the hardware side, the cameras, LED indicators, and Jetson are encapsulated inside the case with only the necessary items accessible from the outside. Because the Flask server is set to run on boot, a user is able to connect and use the GUI without interacting with the Jetson through SSH or connecting peripherals. Lastly, the GUI was created to be simple and intuitive for everyday use and features the use of Pannellum to pan around the stitched video.

To that end, user testing was conducted to improve the GUI. Surveyed individuals were observed using the GUI and asked to give feedback. Some resulting improvements include: improved spacing, choosing different icons, and adding tooltips to buttons.

VII. Impact

The primary contribution of the custom stitcher in our camera setup is increased availability of alternative open-sourced stitching algorithms. While it is available for general use, it also offers developers a robust performance testing framework that can be used across different stitchers. These make this project a good candidate for being built upon with other systems like object detection. It also allows developers to experiment and collect data with other stitchers, processors, cameras, etc. While OpenCV has a default stitcher that implements the full stitching pipeline above, our stitcher has proper stitching functionality in 360-degree configurations through our cylindrical stitching implementation. Furthermore, due to the fixed camera position design, the custom stitcher has better performance in fps rate and reduced stitching error rate.

Our current implementation comes with drawbacks relating primarily to stitch quality and suboptimal framerate. Specifically, compositing effects that are applied to improve the image quality either were unable to be precomputed or caused a significant performance drop. With these limitations,

it would be unsuitable for security applications where identifying an individual and timing might be important. The ideal scenario for this project would be something like an inventory tracker explained in the introduction. In this case, stitching algorithm would work as a sort of preprocessing step for multiple camera streams that utilizes the full scene and prevents objects from being duplicately counted.

For future improvement, the major holdback to the efficiency of the program is the slow frame rate relative to cameras in everyday use. The image stitching step has the longest delay; optimization of this phase of processing the images would greatly improve our stitcher's ability to be applicable to real-world cases where a frame rate of around 30 is usually acceptable. Due to this, any additional systems would likely require more computationally powerful hardware.

VIII. Conclusion

This report covered the computer vision concepts required in this project. Combining this with the results from Alpha and Beta testing, it is shown why the decision was made to create a custom stitcher for this project.

In this production release of the 360° camera, all basic functionality including the ability to view & stitch from all cameras has been implemented. With the pending addition of cylindrical warping, the image coherency can be improved and provide a full 360° image instead of a dual 180°-stitched display output. This will deliver on the overall goal of this project, with future goals focused on image quality through improved stitching & compositing. This makes this project suitable to be built upon with future projects as discussed in the Introduction.

IX. References

- [1] Microchip, "USB5807" *Microchip Technology Inc.* [Online]. Available: <https://www.microchip.com/en-us/product/USB5807> [Accessed: 28-Feb-2022].
- [2] Microchip, "USB2517" *Microchip Technology Inc.* [Online]. Available: <https://www.microchip.com/en-us/product/USB2517> [Accessed: 02-Mar-2022].
- [3] Microchip, "EVB-USB5807" *Microchip Technology Inc.* [Online]. Available: <https://www.microchip.com/en-us/development-tool/EVB-USB5807> [Accessed: 02-Mar-2022].
- [4] Microchip, "EVB-USB2517" *Microchip Technology Inc.* [Online]. Available: <https://www.microchip.com/en-us/development-tool/EVB-USB2517> [Accessed: 02-Mar-2022].
- [5] Findchips, "USB2517" *Findchips.com* [Online]. Available: <https://www.findchips.com/search/USB2517> [Accessed: 03-Mar-2022].
- [6] Findchips, "USB5807" *Findchips.com* [Online]. Available: <https://www.findchips.com/search/USB5807> [Accessed: 03-Mar-2022].
- [7] Electronicsnotes, "USB Standards: USB 1, USB 2, USB 3, USB 4 - capabilities & comparisons" *Electronicsnotes.com*. [Online]. Available:

<https://www.electronics-notes.com/articles/connectivity/usb-universal-serial-bus/standards.php> [Accessed: 01-Mar-2022].

[8] Microchip Support, “Cascading Hubs” *Microchip Technology Inc.* [Online]. Available: <https://microchipsupport.force.com/s/article/USB-HUBS---Cascading-hubs> [Accessed: 03-Mar-2022].

[9] J. R. Picture and Jay Rambhia Senior Android Developer Also, “Stereo calibration,” *Jay Rambhia's Blog*, 30-Mar-2013. [Online]. Available: <https://jayrambhia.com/blog/stereo-calibration>. [Accessed: 04-Apr-2022].

[10] Harvard University, “Real-time image stitching,” *CS205 Real Time Image Stitching*, 2018. [Online]. Available: <https://cs205-stitching.github.io/>. [Accessed: 07-Apr-2022].

[11] Nghonda, Erman, “Pose Estimation and Triangulation” *University of Florida, Event: Research*. [Online]. Available: File request. [Accessed: 07-Apr-2022].

[12] Gray, Dan and Lee, Keven, “360 Camera” *Hackaday.IO*. [Online]. Available: <https://hackaday.io/project/11604-360-camera>. [Accessed: 07-Apr-2022].

[13] Rosebrock, Adrian, “Image Stitching with OpenCV and Python” *pyimagesearch*. [Online]. Available: <https://pyimagesearch.com/2018/12/17/image-stitching-with-opencv-and-python/> [Accessed: 07-Apr-2022].

[14] OpenCV, “Camera calibration,” OpenCV. [Online]. Available: https://docs.opencv.org/3.4/dc/dbb/tutorial_py_calibration.html. [Accessed: 07-Apr-2022].

[15] “Basic concepts of the homography,” *OpenCV*. [Online]. Available: https://docs.opencv.org/4.x/d9/dab/tutorial_homography.html. [Accessed: 27-Oct-2022].